

Anshik Mantri

Data Scientist | Machine Learning Engineer | Backend Developer
+91-7023805961 | anshikmantri26@gmail.com | [LinkedIn](#) | [GitHub](#) | Bengaluru, India

TECHNICAL SKILLS

Languages: Python, SQL, Java, R, Shell Scripting, C++

ML / Deep Learning: Scikit-learn, PyTorch, TensorFlow, Keras, XGBoost, Random Forest, Transformers, NLTK, Hugging Face

AI / NLP / RAG: LangChain, FAISS, RAG Pipelines, Sentence Transformers, Multi-Agent Systems, ChromaDB, Ollama, Vector Search

Big Data: Apache Spark, Hadoop (MapReduce), Hive, ETL Pipelines, Pandas, NumPy

Databases: MySQL, PostgreSQL, MariaDB, MongoDB, Redis, Aerospike, Neo4j — relational, non-relational, graph

Visualization: Power BI, Plotly, Altair, Streamlit, Matplotlib, Seaborn

Cloud & DevOps: AWS (Certified Cloud Practitioner), Microsoft Azure, Docker, Kubernetes, CI/CD, Jenkins, Git, DataDog, Linux/Unix

Backend: FastAPI, Spring Boot, REST APIs, Microservices, JWT, OAuth2, API Gateway, Rate Limiting

PROFESSIONAL EXPERIENCE

Backend Developer Intern | PharmEasy | Bengaluru, India

Sep 2025 – Present

- **B2B API Suite:** Developed and enhanced B2B backend APIs (orders, reports, authentication, catalog) with validation, filtering, pagination, and guard-based authentication, enabling scalable partner integrations across distributed services.
- **Rate Limiting & Request Locking:** Implemented Redis-based rate limiting and concurrent request locking for the Products API, eliminating duplicate requests and enforcing per-minute API usage quotas; resolved production bugs by restructuring controller flow for safe Redis key lifecycle management.
- **Security (IDOR Fix):** Eliminated a critical IDOR vulnerability in order management APIs by enforcing JWT-based ownership validation and HTTP 403 access control, blocking unauthorized access and aligning with OWASP security best practices.

AI / ML Engineer | Rofolin | Remote

Aug 2025 – Sep 2025

- **Multi-Agent Platform:** Built a Python-based AI automation system with specialized LLM agents (research, trend analysis, content generation) orchestrated via a central controller, automating end-to-end data-driven marketing campaigns.
- **RAG Pipeline:** Engineered a production-grade RAG pipeline using FAISS vector search and Sentence Transformers, backed by a FastAPI service and real-time analytics dashboard; enabled context-aware LLM content generation at scale on Linux cloud infrastructure.

Data Analyst Intern | HexaWealth | Remote

Aug 2024 – Jan 2025

- **Data Engineering & EDA:** Built ETL pipelines to ingest multi-source financial data into MySQL warehouses; performed EDA on NIFTY/BSE 500 equities (Pandas, NumPy) covering sector returns, volatility trends, and drawdown patterns across 500+ instruments.
- **Statistical Analysis & Risk:** Developed a risk classification engine using SQL window functions and Python to segment assets by Sharpe ratio, volatility, and return quartiles; validated signals using t-tests and chi-square tests for statistical significance.
- **Predictive Modeling:** Trained and evaluated Random Forest, XGBoost, and Logistic Regression models on historical equity and macroeconomic data using Scikit-learn pipelines with feature engineering, cross-validation, and hyperparameter tuning.

KEY PROJECTS

StockSense AI — Real-Time Stock Analysis Platform

Python | MySQL | Scikit-learn | XGBoost | LangChain | Gemini LLM | NLP | Plotly | Altair

- **Integrated LangChain and Gemini LLM** as a finance-domain AI assistant with TextBlob sentiment analysis on live news feeds and a custom screener filtering by P/E, ROE, and Debt/Equity; backed by Random Forest and XGBoost pipelines with feature engineering and cross-validation on NIFTY/BSE 500 instruments.
- **Built an AI-driven portfolio builder** with dynamic risk profiling, SIP/lump-sum projections, and multi-stock comparison surfaced through Plotly/Altair dashboards tracking volatility, drawdown, sector allocation, and benchmark deviation in real time.

Local-Vault — Privacy-First AI Chatbot

Python | RAG | ChromaDB | Ollama | FastAPI | Streamlit

- **Built a full-stack AI chatbot** with a RAG pipeline supporting document ingestion (PDF, Word, CSV, images) using OCR, semantic chunking, metadata tagging, and ChromaDB vector search to enable accurate, context-aware LLM responses.
- **Developed a Streamlit interface** with FastAPI backend and integrated Ollama for local LLM inference, enabling private on-device AI, conversation memory, and personalized responses from user-uploaded data.

EDUCATION

B.Tech, Computer Science Engineering | Vellore Institute of Technology, Chennai

2021 – 2025

GPA: 8.1 / 10.0 | Coursework: DSA, DBMS, Statistics & Probability, Data Visualization, Machine Learning, Distributed Systems, Computer Networks, OS, OOP

CERTIFICATIONS & ACHIEVEMENTS

- AWS Certified Cloud Practitioner — Amazon Web Services
- Data Science, ML, Deep Learning & NLP Bootcamp — Udemy | Data Engineering, Python for Data Science & AI, Apache Spark & Hadoop — IBM / Coursera
- Mindcare Hackathon — Top 5 Finalist (200+ teams) | Amazon ML Hackathon — Participant